

# EvolveDB - Documentation for End-User

*Author:* Torben Eckwert, M.Sc.  
*E-Mail:* torben.eckwert@zdh.thm.de

*Subject Area:* Computer Science

*Supervisors:* Prof. Dr. rer. nat. Michael Guckert  
Prof. Dr. ing. Gabriele Taentzer

December 8, 2022 , Wetzlar



Philipps



Universität  
Marburg



---

## FORSCHUNGSCAMPUS MITTELHESSEN

# Contents

1	Introduction	1
2	Prerequisites and Installation	1
3	EvolveDB	3
3.1	Reverse Engineering	4
3.2	Restructuring	8
3.2.1	Change Column Type	9
3.2.2	Add New Columns	9
3.2.3	Change Default Values	11
3.2.4	Reduce Column Size	12
3.2.5	Change the Multiplicity of an Association	12
3.3	Model Comparison	14
3.3.1	SiLift	14
3.3.2	Model Comparison with EvolveDB	16
3.4	Forward Engineering	21
3.4.1	Migration Model	21
3.4.2	Breaking & Resolvable	24
3.4.3	Not Automatically Resolvable	26
3.5	Generate Migration Scripts	27
4	Migration Scripts	29
4.1	History Table	33

# 1 Introduction

EvolveDB is an Eclipse-based framework for schema evolution in MySQL databases. The user specifies the evolution steps by freely editing a database model extracted by reverse engineering. EvolveDB analyzes the differences between the status quo and the evolved model structures and generates a data migration script. This user manual includes installation instructions and introductory tutorials on how to use EvolveDB as an end user.

## 2 Prerequisites and Installation

EvolveDB is a plug-in for recent versions of the [Eclipse Modeling Tools](#) (last tested version: 2022-06). The Eclipse feature is available from our [GitHub repository](#).

**Important: Currently, EvolveDB requires Java SE 11. Newer or older Java Versions may cause unexpected behavior. The Java version can be changed by modifying the vm option in the *eclipse.ini*.**

Before installing EvolveDB, we need to install some additional Eclipse plugins. The following list contains the mandatory plugins.

- [Sirius](#) is an Eclipse project which allows you to easily create your own graphical modeling workbench by leveraging the Eclipse Modeling technologies, including EMF and GMF.
- [Eclipse OCL](#) is available via the eclipse marketplace.

If all requirements are fulfilled, EvolveDB can be installed via the menu item **Help** ⇒ **Install New Software...** (fig. 1).

We used the [MySQL Server](#) version 8.0.28 for this tutorial. When using older versions, compatibility problems may occur.

## 2 Prerequisites and Installation

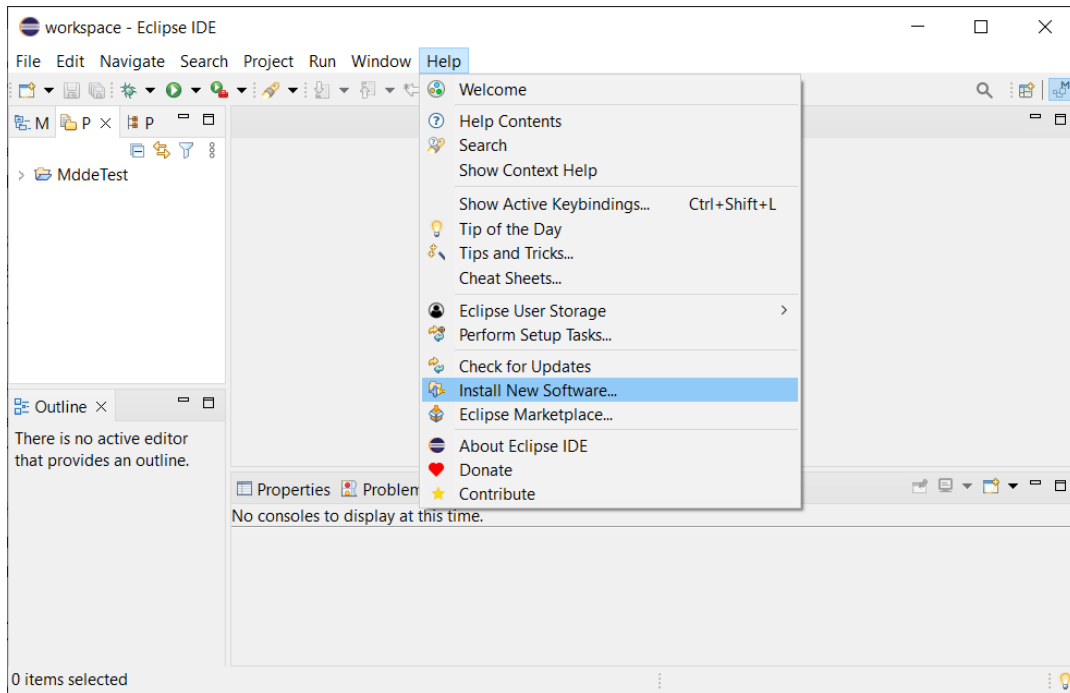


Figure 1: Eclipse: Install New Software...

The repository contains multiple categories (fig. 2). For the following tutorial, all elements must be installed.

### 3 EvolveDB




Figure 2: Eclipse: EvolveDB local repository

## 3 EvolveDB

The schema evolution with EvolveDB is a model-driven reengineering process that includes three phases, reverse engineering, restructuring, and forward engineering. In the first step, we start by creating a representation of a (MySQL) database schema through reverse engineering. In the restructuring phase we edit this model resulting in a second model version. Afterwards EvolveDB analyzes the differences between the status quo and the evolved model structures. From the delta between the two model versions, SQL migration scripts are generated, which can be used to migrate the schema and the associated data.

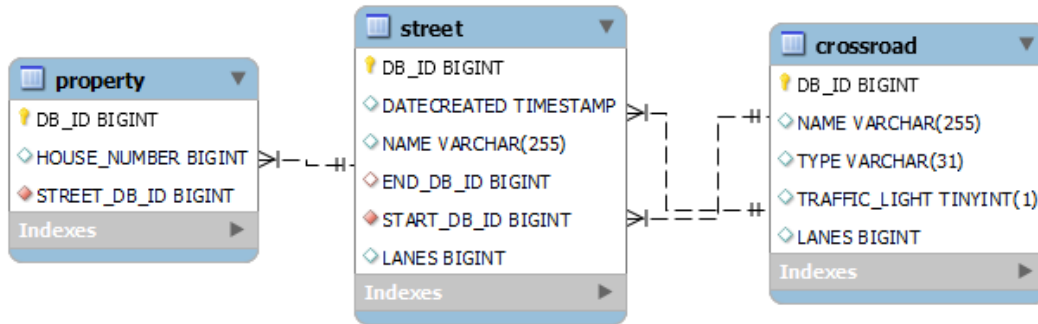


Figure 3: MySQL database schema

The use of EvolveDB as a tool for schema evolution is demonstrated with a running example. The starting point is an already existing MySQL database. (fig. 3) The schema describes traffic routes (table *street*). *Streets* always have a starting point. These starting points are intersections or traffic circles (table *crossroad*). Several roads can start or end at one node. Furthermore, properties (table *property*) can be located at traffic routes. During this tutorial, we are going to make the following changes to the database schema and the associated data:

- In the table *street* we change the type of the column *DATECREATED* changed from *TIMESTAMP* to *DATETIME*.
- We add two new columns, *LONGITUDE* : *Integer* and *LATITUDE* : *Integer*, to table *crossroad*.
- We reduce the column size of the *NAME* column in table *street* from 255 to 40.
- We change the single-valued association between property and street into a multi-valued association. (Requires a new cross-reference table.)
- We set 1 as new default value for column *lanes* in table *crossroad*.

### 3.1 Reverse Engineering

In the first step we extract a database model by reverse engineering. For this purpose, open the Eclipse Project Explorer and create a new project. Then select the new project and open the context menu with a right click. Select **EvolveDB** ⇒ **Create MDDE Model...** (Fig. 4)

### 3 EvolveDB

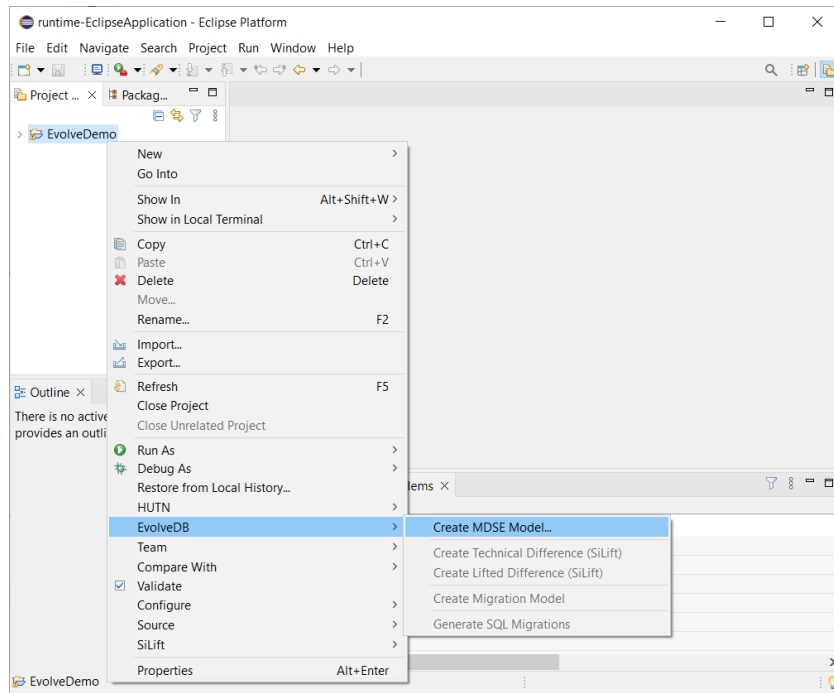


Figure 4: Start reverse engineering

A wizard dialog opens that has several pages. On the first page, we have to select the data source. (fig. 5)

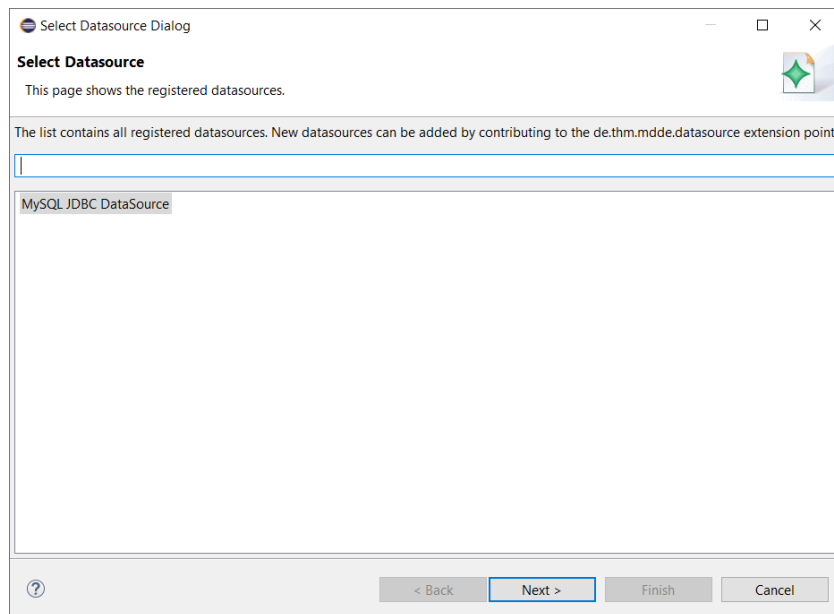


Figure 5: Select data source

EvolveDB does not include database drivers. When used for the first time, we must down-

load the JDBC driver. The driver files can be downloaded automatically. Or you can obtain driver files by yourself and add them to the driver directory. The default driver location is the eclipse install directory. The location can be changed in the eclipse preferences. If no suitable driver can be found, EvolveDB opens a download driver dialog. (fig. 6)

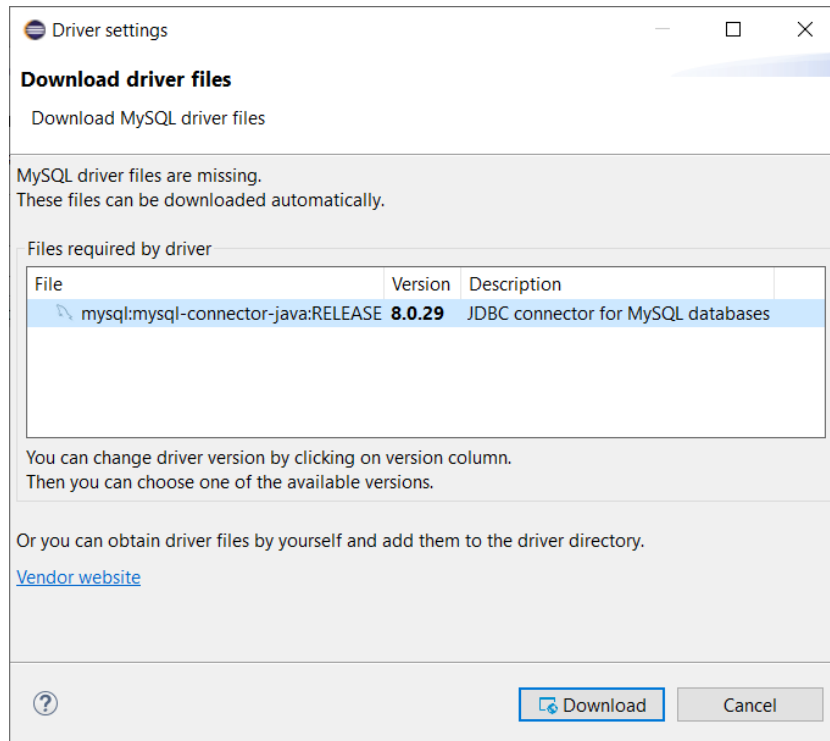
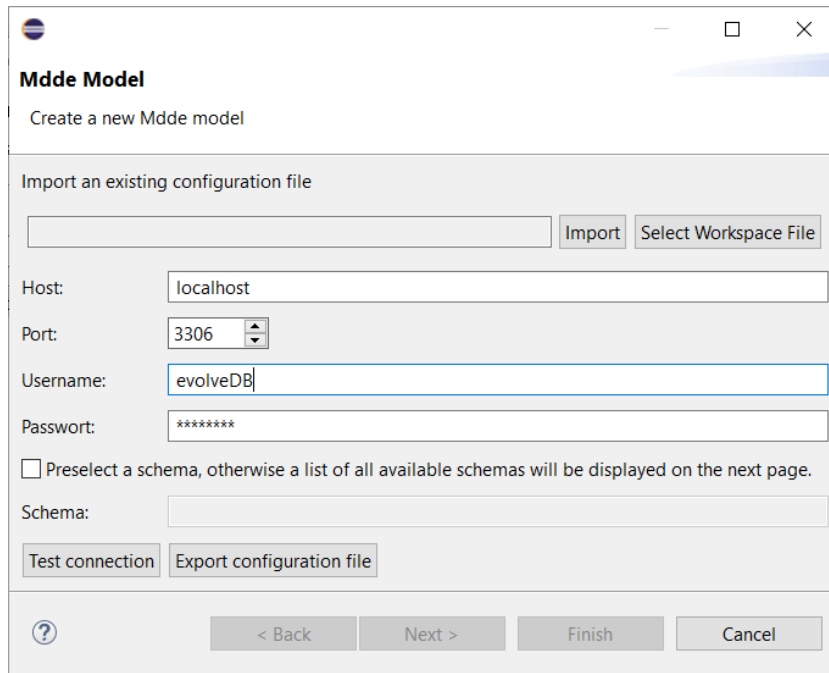


Figure 6: Download JDBC driver

After loading the driver class, we have to insert the connection information. (fig. 7) Furthermore, it is possible to save the connection information (without password) as an XML configuration file. This configuration file can be imported when creating a new model so that the connection information does not have to be entered again. (fig. 9) After entering the required information, the connection must be tested. We can continue if the connection is successful. The next page shows a list with all available database schemes if we did not preselect a schema on the first page. (fig. 8) On the last page of the dialog, the new model's name and location must be specified. The dialog can then be closed with *Finish*, and the newly created model is opened in the editor. In addition, a new folder with the name *genModel* and a copy of the model are created in the storage location. The second model is required for the comparison process after the restructuring phase and should not be changed.



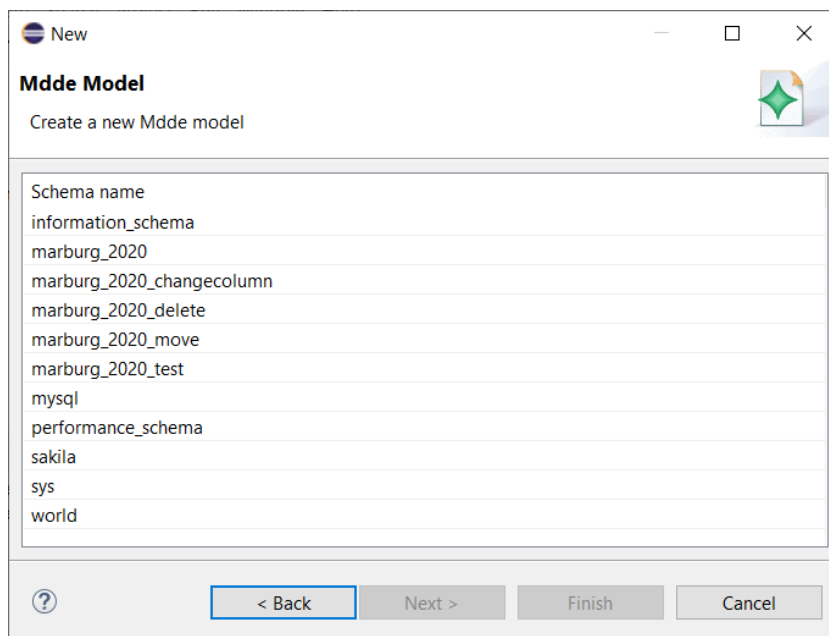
### 3 EvolveDB



The screenshot shows a dialog box titled "Mdde Model" with the subtitle "Create a new Mdde model". It contains the following elements:

- An "Import an existing configuration file" section with a text input field, an "Import" button, and a "Select Workspace File" button.
- Fields for "Host" (localhost), "Port" (3306), "Username" (evolveDB), and "Password" (masked with asterisks).
- A checkbox labeled "Preselect a schema, otherwise a list of all available schemas will be displayed on the next page." which is currently unchecked.
- A "Schema:" text input field.
- "Test connection" and "Export configuration file" buttons.
- Navigation buttons: "?", "< Back", "Next >", "Finish", and "Cancel".

Figure 7: Insert connection data



The screenshot shows the same "Mdde Model" dialog box, but now displaying a list of available database schemas. The "Preselect a schema" checkbox is checked. The list of schemas includes:

- information\_schema
- marburg\_2020
- marburg\_2020\_changecolumn
- marburg\_2020\_delete
- marburg\_2020\_move
- marburg\_2020\_test
- mysql
- performance\_schema
- sakila
- sys
- world

The "Next >" button is highlighted, indicating the user is moving to the next step in the wizard.

Figure 8: Select database schema

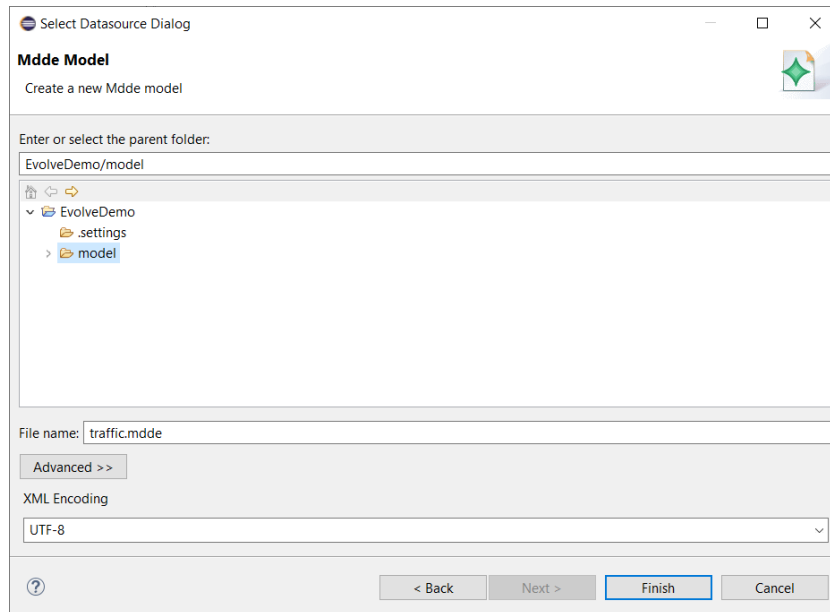


Figure 9: Choose file location

## 3.2 Restructuring

In the previous step, we created a database model by reverse engineering. The new model was automatically opened in the Eclipse editor (fig. 10). The editor consists of two parts. On the left is a tree-based editor that allows navigating in the model, and on the right is a view with the selected element's properties. Each element in our model corresponds to an element in the database. By expanding an element, the child elements become visible. Now we can start editing the model.

### 3 EvolveDB

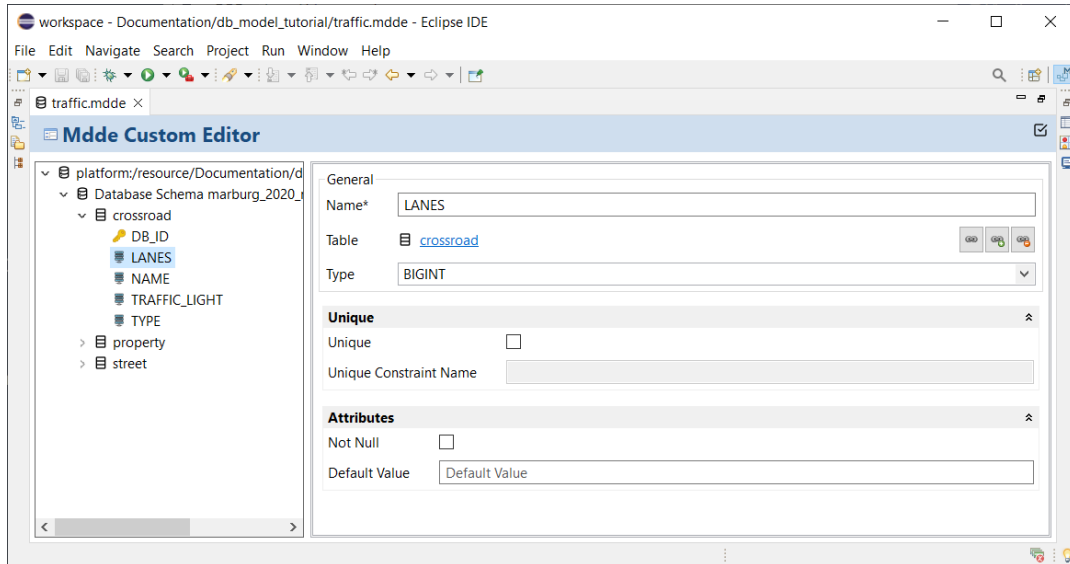


Figure 10: MDDE model editor

#### 3.2.1 Change Column Type

Select the column *DATECREATED* in table *street*. The attributes of the column are now displayed in the right side of the editor. Change the type from *TIMESTAMP* to *DATETIME*.

#### 3.2.2 Add New Columns

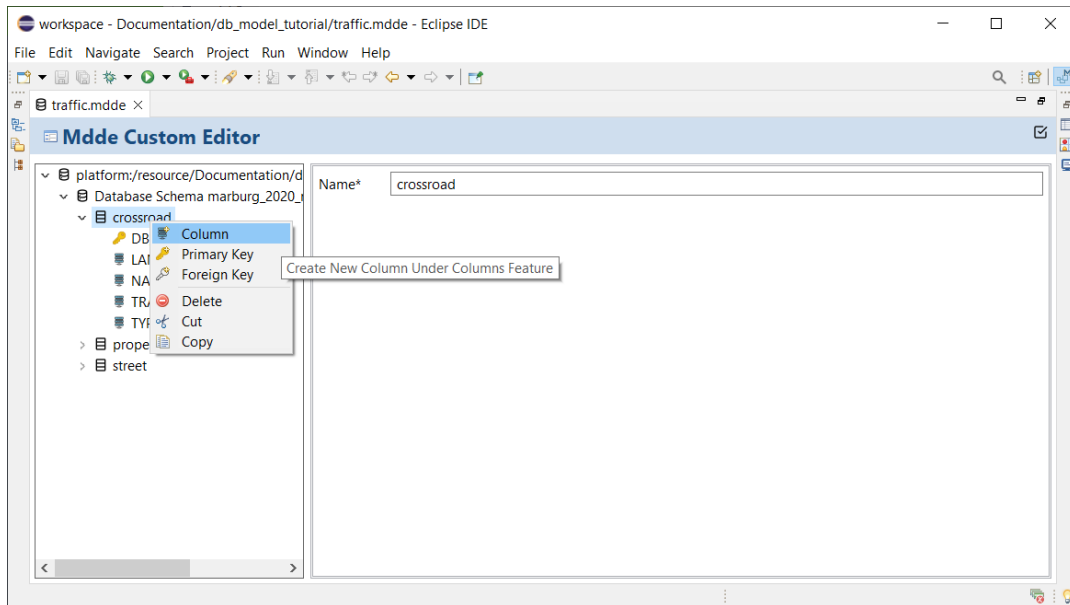


Figure 11: Context menu

### 3 EvolveDB

In this step we are going to add the two new columns for longitude and latitude. Select the *crossroad* table and call the context menu with a right-click (fig. 11). Create a new *column*. The new column will then be opened (fig. 12). Mandatory fields are marked in red. Name the new attribute *LATITUDE* and set the type to *Integer*. Since columns with the *Integer* type do not have a *size* attribute, the attribute is hidden. The other attributes can remain unchanged. (fig. 13) Repeat the process for the *LONGITUDE* column.

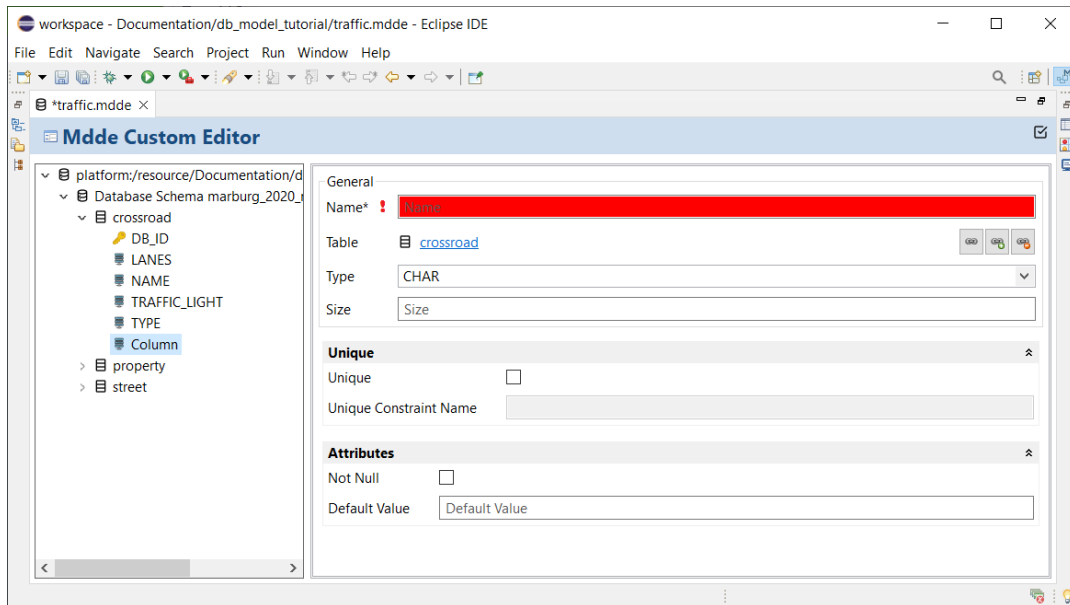


Figure 12: Mandatory fields are marked in red

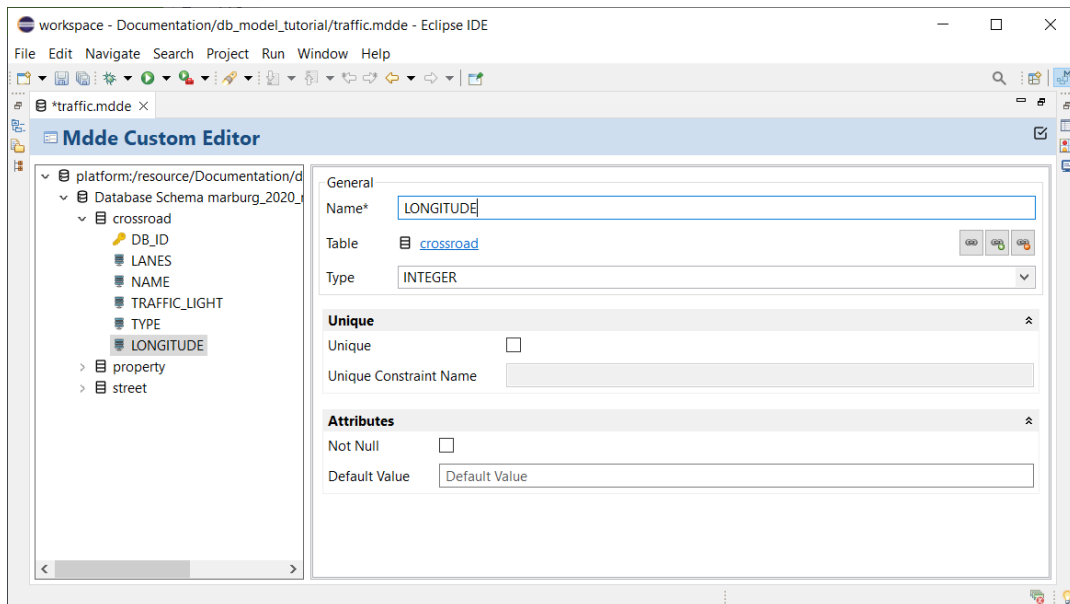


Figure 13: The *Size* attribute is not visible

### 3.2.3 Change Default Values

Open the *crossroad* table in the tree viewer and select the column *LANES*. Set the default value to 1. Basically, this change is no different from a name change or a type change. Therefore, you may wonder at this point why this change is listed and explained in this tutorial. The reason is one of the significant strengths of EvolveDB. We explain the details in the following section.

**Explanation: Default-Value, Size and Validation** Each column in a MySQL database can have a default value. However, this default value must match the column type. If we consider the column *LANES* as an example, the default value may only consist of digits because the type is a numeric data type. However, this is not the only restriction. At the same time, the default value must not exceed the value range of the data type. The *Integer* data type has a fixed size. Therefore, EvolveDB hides the size attribute for columns with the *Integer* data type. Columns of type *Integer* can contain values from the range between -2147483648 and 2147483647. When we specify a value outside the value range or enter invalid characters, we receive an error message when validating the model (fig. 14). The validation can be triggered manually using the button in the top right corner of the editor.

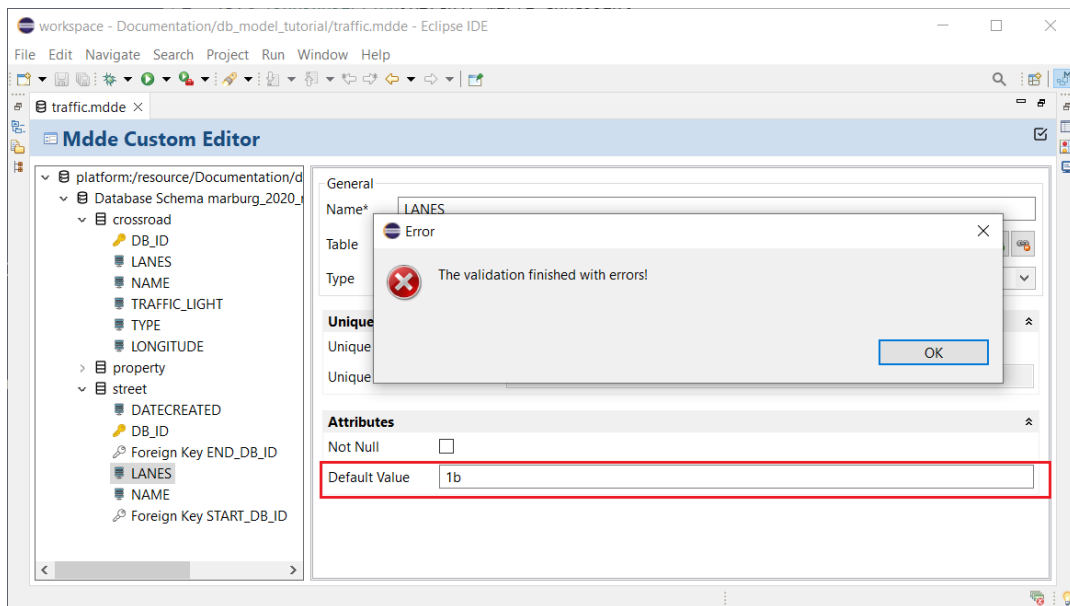


Figure 14: The default value is invalid.




Figure 15: Default-value for a timestamp column

Let us consider another example. MySQL allows fractional seconds for the *TIME*, *DATE-TIME* and *TIMESTAMP* data types with an accuracy of up to microseconds (6 digits). When validating the default value of a column with one of these three types, the default value must match the *Size* attribute (fig. 15). In summary, a simple change to the default value or size attribute requires extensive validation. These validations ensure that only permissible values can be entered.

### 3.2.4 Reduce Column Size

Select the column *NAME* in table *street*. The column has the datatype *VARCHAR*. Unlike numeric data types like *INTEGER*, *VARCHAR* has no fixed size. For this reason, the size attribute is visible, and we can select a value in the range between 0 and 255. Currently, 255 is selected. Decrease the value to 40.

### 3.2.5 Change the Multiplicity of an Association

Right-clicking on the database schema within the tree-based editor opens the context menu. Via the context menu, we can add new tables (fig. 16). In the previous sections, the extensive validations have already been discussed. However, not only individual columns or attribute values are validated, but also entire tables and the model as a whole. After we add the new table *property\_street*, the validation of the model leads to several errors (fig. 17). The model is invalid, because we have not added any columns yet.

### 3 EvolveDB

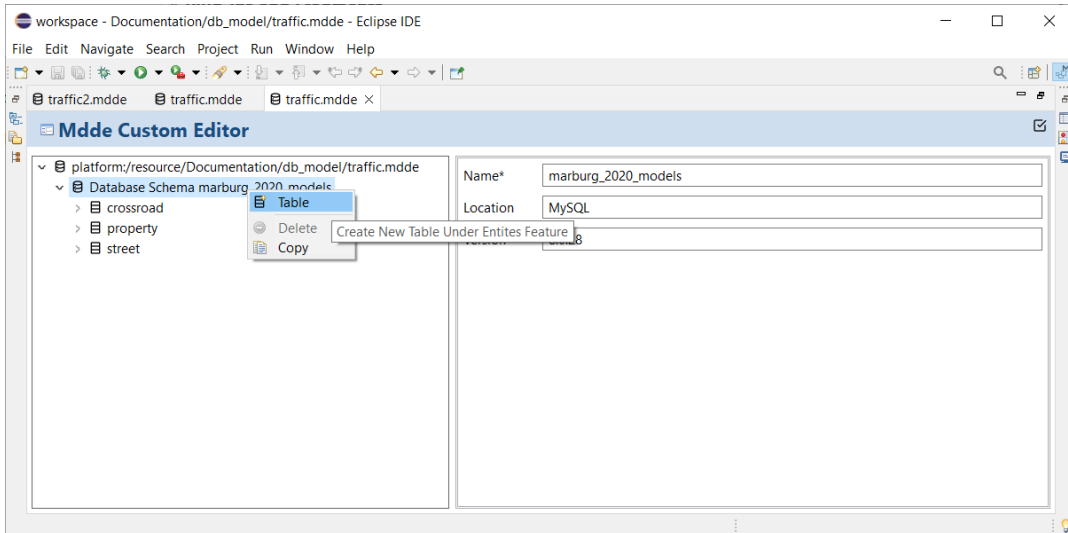


Figure 16: Add new table

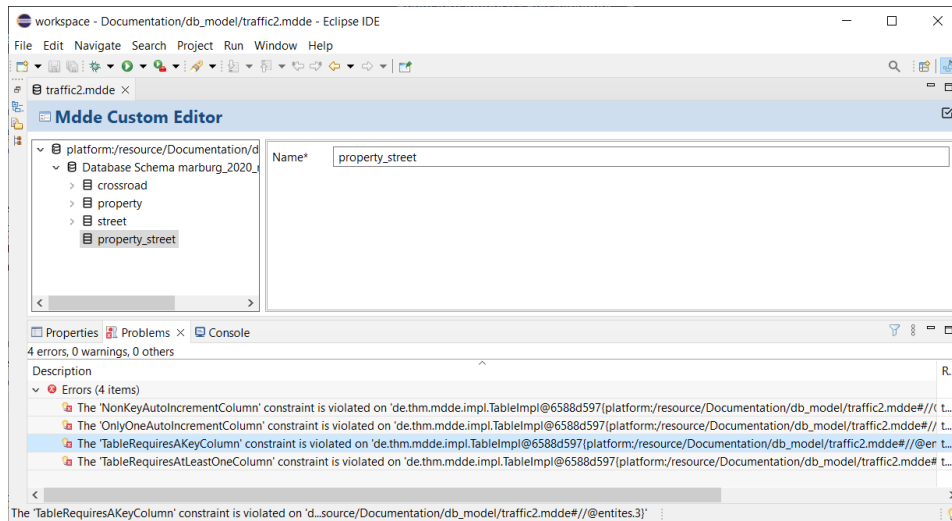
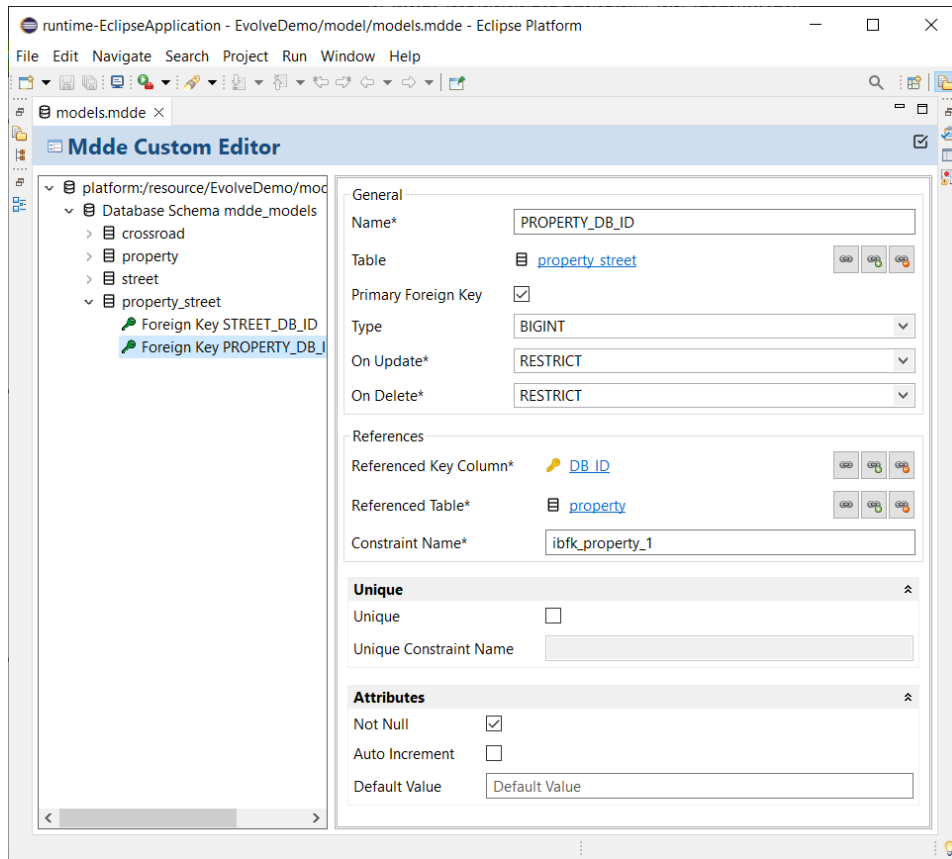


Figure 17: Model validation

Tables must have at least one column and a primary key. By double-clicking on the error, the model element that caused the error is selected. Now we add the two Foreign Keys to the table. The two foreign keys reference the tables *street* and *property*. (It is also possible to move the *STREET\_DB\_ID* to table *property\_street*.) Then the old foreign key *street\_db\_id* in the table *property* can be deleted. Figure 18 shows the the result. Once the validation of the model is successful, the comparison of the model versions can be continued in the next step.

Figure 18: Table *property\_street* with foreign keys

### 3.3 Model Comparison

EvolveDB analyzes the differences between the two model versions. At first, a sequence of atomic differences is automatically derived. Then the identified model evolution steps are (semi-)automatically mapped to predefined migration operations. These migration operations can finally be applied to the original database schema to transform it into the target schema. For the comparison process, EvolveDB uses the *SiLift* framework. *SiLift* is a generic model comparison environment that can semantically lift identified low-level differences of EMF-based models into representations of user-level edit operations.

#### 3.3.1 SiLift

SiLift's approach can best be compared to a four-step pipeline, as shown in Figure 19. The two model versions, e.g. *models.mdde* and *models2.mdde* serve as input:

1. **Matching:** The task of a matcher is to identify the corresponding elements from model A and model B, i.e., the elements that match in both models. SiLift provides four different matchers by default. One of them is EMFCompare.






Figure 19: SiLift processing pipeline

2. **Difference Derivation:** Based on the correspondences, the difference derivator calculates a technical difference (low-level difference) between the model versions. All objects and references for which no correspondence exists must either have been added to model B or removed from model A.

```

Δ platform:/resource/Documentation/db_model_tutorial/test.symmetric:
  Δ Symmetric-Difference: trafficV2.mdde <-> traffic.mdde
    ✖ Remove-Object: STREET_DB_ID
    ✖ Remove-Reference: table (STREET_DB_ID -> property)
    ✖ Remove-Reference: referencedTable (STREET_DB_ID -> street)
    ✖ Remove-Reference: referencedKeyColumn (STREET_DB_ID -> DB_ID)
    + Add-Object: LATITUDE
    + Add-Reference: table (LATITUDE -> crossroad)
    + Add-Object: LONGITUDE
    + Add-Reference: table (LONGITUDE -> crossroad)
    + Add-Object: property_street
    + Add-Reference: schema (property_street -> marburg_2020_complex)
    + Add-Reference: columns (property_street -> STREET_DB_ID)
    + Add-Reference: columns (property_street -> PROPERTY_DB_ID)
    + Add-Object: STREET_DB_ID
    + Add-Reference: table (STREET_DB_ID -> property_street)
    + Add-Reference: referencedTable (STREET_DB_ID -> street)
    + Add-Reference: referencedKeyColumn (STREET_DB_ID -> DB_ID)
    + Add-Object: PROPERTY_DB_ID
    + Add-Reference: table (PROPERTY_DB_ID -> property_street)
    + Add-Reference: referencedTable (PROPERTY_DB_ID -> property)
    + Add-Reference: referencedKeyColumn (PROPERTY_DB_ID -> DB_ID)
    + Add-Reference: entites (marburg_2020_complex -> property_street)
    Δ Attribute-Value-Change: DATECREATED -> DATECREATED
    Δ Attribute-Value-Change: NAME -> NAME
    ✖ Remove-Reference: referencedBy (DB_ID -> STREET_DB_ID)
    + Add-Reference: referencedBy (DB_ID -> STREET_DB_ID)
    + Add-Reference: columns (crossroad -> LATITUDE)
    + Add-Reference: columns (crossroad -> LONGITUDE)
    Δ Attribute-Value-Change: LANES -> LANES
    ✖ Remove-Reference: columns (property -> STREET_DB_ID)
    + Add-Reference: referencedBy (DB_ID -> PROPERTY_DB_ID)
  > Matching (trafficV2.mdde <-> traffic.mdde)
  > platform:/resource/Documentation/db_model_tutorial/genModel/trafficV2.mdde
  > platform:/resource/Documentation/db_model_tutorial/traffic.mdde
  
```

Figure 20: Technical difference between models.mdde and models2.mdde

3. **Semantic Lifting:** The previously calculated technical difference contains all changes between the models. Then these changes are semantically lifted. For this purpose, the low-level changes are grouped into semantic change sets with the help of recognition rules. Each semantic change set represents an editing operation performed by the user. Often, user editing operations consist of more than one low-level change, because even basic edit operations that appear simple from a users point of view may lead to many low-level changes. While atomic rules cover the creation, deletion, moving of elements, and the changing of attribute values, the complex editing rules are usually composed of atomic and other complex rules.
4. **Difference Presentation UI:** The result is a symmetric difference model which contains the matching, the atomic changes and the resulting semantic change sets.

### 3.3.2 Model Comparison with EvolveDB




Figure 21: Start comparison

We start comparing the two model versions from the previous chapter. Therefore, we select the two **mdde** files in the Package or Project Explorer and open the context menu with a right click. We choose **EvolveDB**  $\Rightarrow$  **Create Lifted Difference (SiLift)** (fig. 21) A multi-page wizard dialog opens. On the first page we must select the model comparison direction. The arrow has to point at the edited model version(fig. 22).

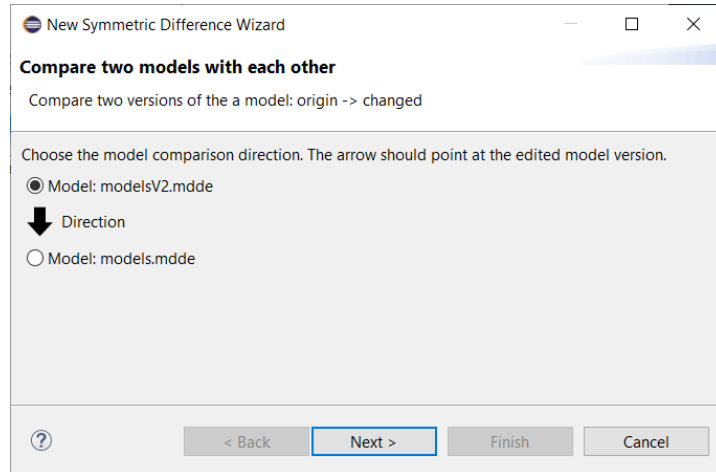


Figure 22: Choose model comparison direction

On the second page (fig. 23) the matching identified by SiLift is presented. The upper table shows the corresponding tables from model A and model B. Unmatched tables which have no correspondence must either have been added to model B (marked in green) or removed from model A (marked in red). The tables at the bottom show the corresponding column. If a column was moved or renamed, the column is marked blue. For example, the foreign key *STREET\_DB\_ID* was moved to the table *property\_street*. If an identified correspondence is not correct, we can delete or add correspondences.

A file name for the symmetric difference model must be specified on the last page of the dialog (fig. 24). Furthermore, we can create a so-called migration model. The migration model will be required later in the process. For this reason, it makes sense to create it directly. The migration model will be stored at the same location as the symmetric difference model.

### 3 EvolveDB

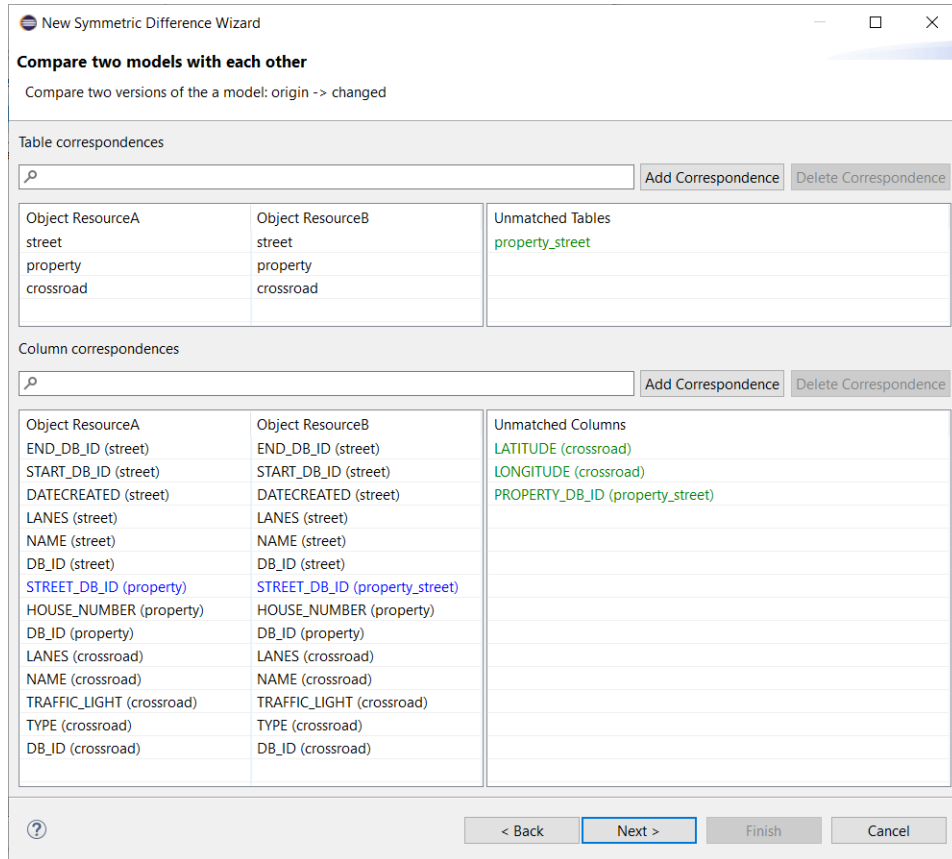


Figure 23: Choose model comparison direction




Figure 24: Choose file location

The symmetric difference model contains the matching, the atomic changes and the resulting semantic change sets. Figure 25 shows the symmetric difference model for our example opened in the symmetric difference model editor. All atomic changes are displayed when we expand one of the semantic change sets. In our example (fig. 25) the complex operator *CHANGE\_1N\_INT0\_NM*, representing the multiplicity change, consists of 21 atomic changes.

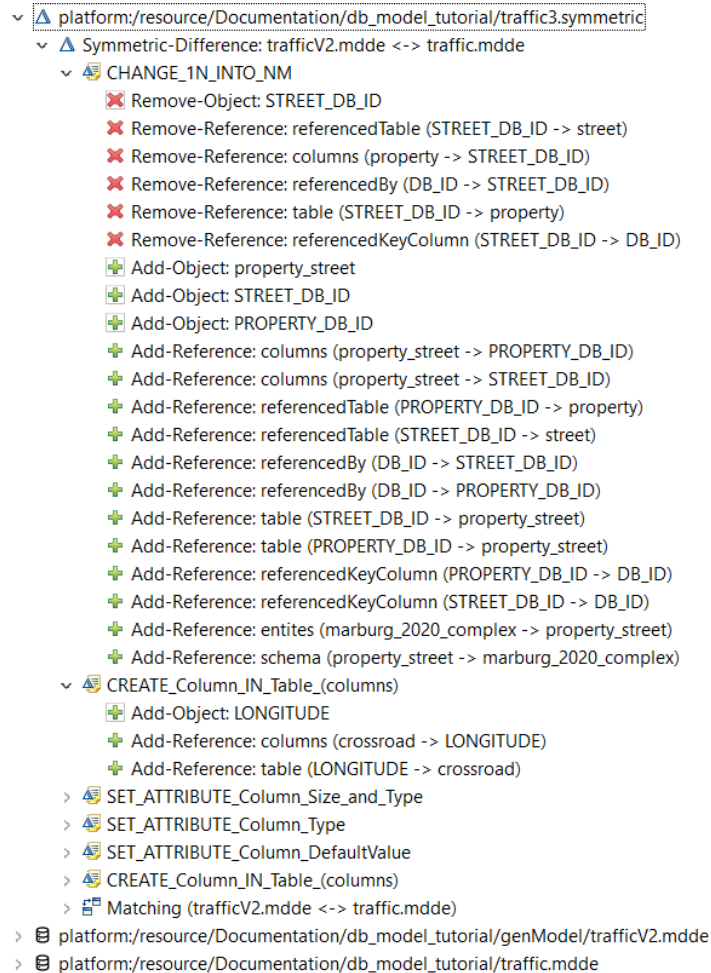


Figure 25: Lifted difference between traffic.mdde and traffic2.mdde

If we recall the changes we did previously, the semantic change sets can be assigned to the editing operations as follows:

- **SET\_ATTRIBUTE\_Column\_Size\_and\_Type**: Reduce the column *size* of *NAME* in table *street* from 255 to 40.
- **CREATE\_Column\_IN\_Table\_(columns)**: Create the new column *LONGITUDE*: in table *crossroad*.

- **CREATE\_Column\_IN\_Table\_(columns)** Create the new column *LATITUDE* in table *crossroad*.
- **CHANGE\_1N\_INTO\_NM**: Change the single-valued association between property and street into a multi-valued association. (Add a new cross-reference table.) This this operator includes both the creation of the new table and the deletion of the old foreign key.
- **SET\_ATTRIBUTE\_Column\_Type**: Change the column type of *DATECREATED* from *TIMESTAMP* to *DATETIME*.
- **SET\_ATTRIBUTE\_Column\_DefaultValue**: Change the default value for column *LANES*.

All currently supported operators are listed in table 1.

Name	Classification
Rename Element	VALUE CHANGE
Make ForeignKey a PrimaryForeignKey	VALUE CHANGE
Set foreign key on update constraint	VALUE CHANGE
Set foreign key on delete constraint	VALUE CHANGE
Set foreign key constraint name	VALUE CHANGE
Make column unique	VALUE CHANGE
Set column unique constraint name	VALUE CHANGE
Set column type	VALUE CHANGE
Set column size	VALUE CHANGE
Set column not null	VALUE CHANGE
Set column default value	VALUE CHANGE
Set column auto-increment attribute	VALUE CHANGE
Drop primary key	REDUCE
Drop table	REDUCE
Drop foreign key	REDUCE
Drop column	REDUCE
Create primary key	INCREASE
Create foreign key	INCREASE
Create table	INCREASE
Create column	INCREASE
Create many-to-many table	STRUCTURAL CHANGE
Move column	STRUCTURAL CHANGE
Dissolve many-to-many table	STRUCTURAL CHANGE
Join table	STRUCTURAL CHANGE

Table 1: Supported edit operations

## 3.4 Forward Engineering

In the next step, we will transform the edit operations into a SQL migration script. An editing operation can be non-breaking, breaking & resolvable, or breaking & unresolvable.

- **non-breaking:** Non-breaking changes can be resolved automatically. Most additions belong to this group.
- **breaking and resolvable:** Breaking and resolvable changes break the conformance of existing data, although they can be automatically adapted.
- **breaking and unresolvable:** Breaking and unresolvable changes break the conformity of existing data which cannot be automatically adjusted and require user intervention. These migrations often require the user to specify how the data migration should be performed. For example, if the column size is reduced, it is necessary to specify how to deal with too large elements.

Changes in the first category are not relevant for the migration problem because data migration is not necessary. Changes placed into the second category can be resolved without user intervention but require data migration. The third category of changes requires user attention. The user has to provide additional information necessary to migrate the affected data.

### 3.4.1 Migration Model

The symmetric difference model we created in the previous section is only used to represent the difference between the two model versions. The difference model is unsuitable for adding the additional information required for the migration. For this reason, we convert the *difference.symmetric* model into a migration model (If this has not yet been done in the previous step). Therefore, we select the model in the Project or Package Explorer and open the menu with a right-click. Then we choose **EvolveDB** ⇒ **Create Migration Model...** (Abb. 26) The migration model is stored in the same folder as the difference model. If the model is not visible immediately, the project or the folder must be refreshed once.

### 3 EvolveDB

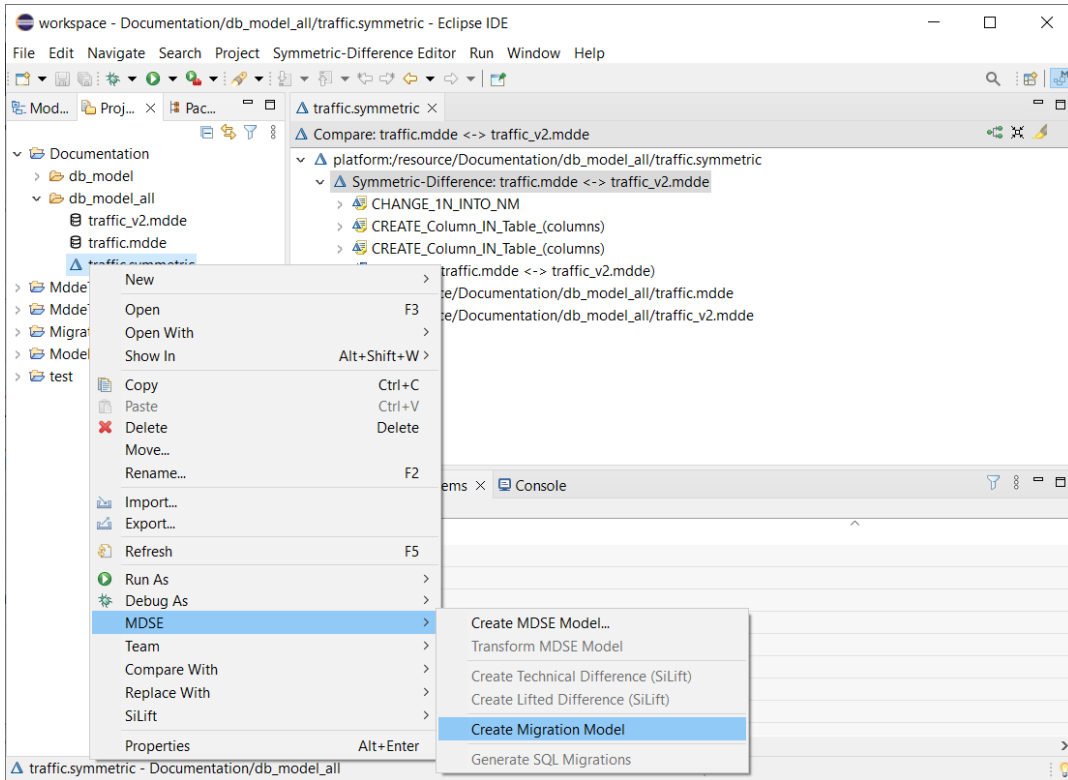


Figure 26: Create migration model

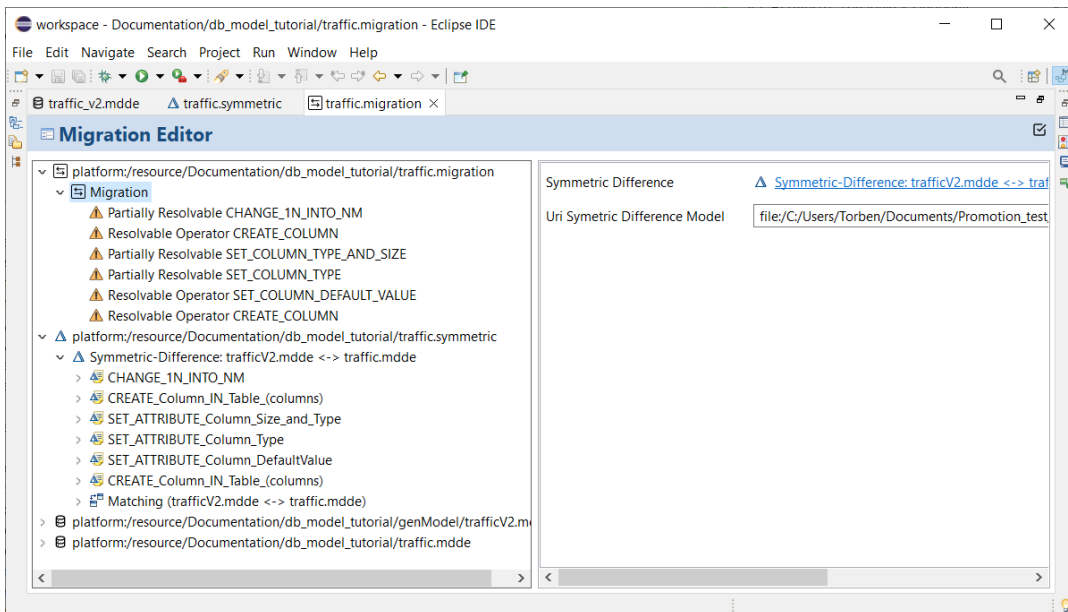


Figure 27: Migration model editor

The new model should be opened automatically in the Migration Editor. This editor also



consists of two areas. On the left side, a tree-based navigation area can be used to navigate through the model. On the right side, the currently selected element is displayed. The migration model references all models created so far, so four models are listed in the navigation view. These four models are the two versions of the traffic model, the difference model, and the migration model (fig. 27).

During the model transformation, the operators were classified according to the classification from the previous section. Operators marked as *Resolvable* belong to the group of *non-breaking* operators. These operators do not affect already existing data. For example, the *CREATE\_COLUMN* operator belongs to this group. Operators marked as *Partially Resolvable* belong to the group of *breaking and resolvable* Operators. Operators of the group *breaking and unresolvable* are prefixed with *Not automatically resolvable*.

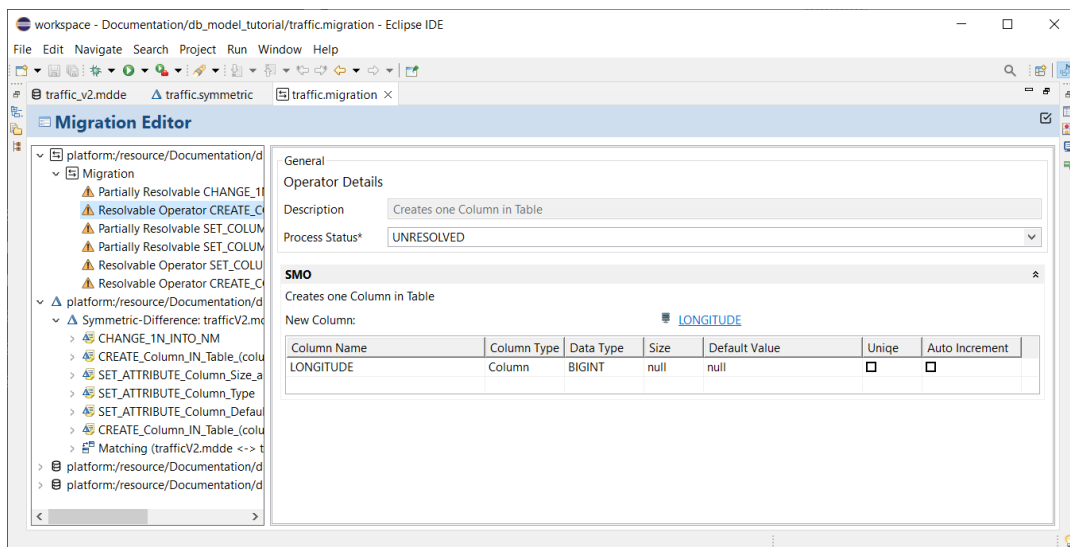


Figure 28: Resolvable operator Create Column

If we select one of the operators, the description opens in the right area of the editor (fig. 28). Every operator should be reviewed by the developer. This precaution may be necessary as not all change operations are covered yet. Therefore, every operator has a so-called process status. We can change the status by selecting a value from the drop-down menu below the operator's description. The three possible values are:

- **UNRESOLVED:** Unresolved is the default value. After creating the migration model, all operators are initially assigned this value. This status means that the operator has not been checked and confirmed by the user yet.
- **RESOLVED:** This status means that the operator has been checked by the user and should be taken into account when subsequently generating the migration scripts.
- **IGNORE:** If an operator is incorrect, it is not necessary to repeat the comparison or the creation of the migration model. Instead, the status can be set to ignore. Operators with this status will not be considered afterwards.

### 3 EvolveDB

All operators with the status *UNRESOLVED*, are marked with a warning in the tree-based editor. As soon as we choose another status, a green checkmark appears instead of the warning (fig. 29).

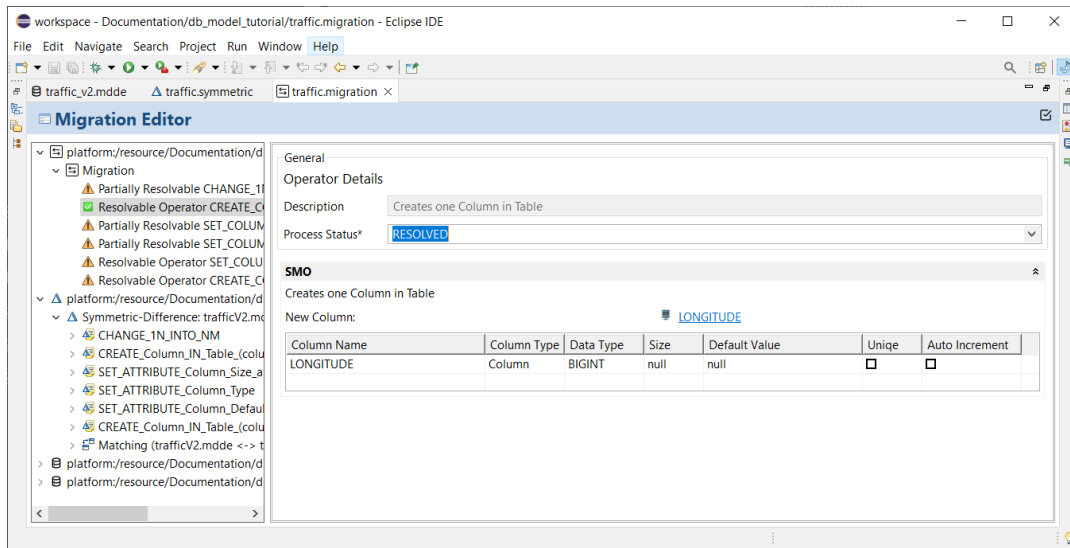


Figure 29: Operator was marked as resolved

#### 3.4.2 Breaking & Resolvable

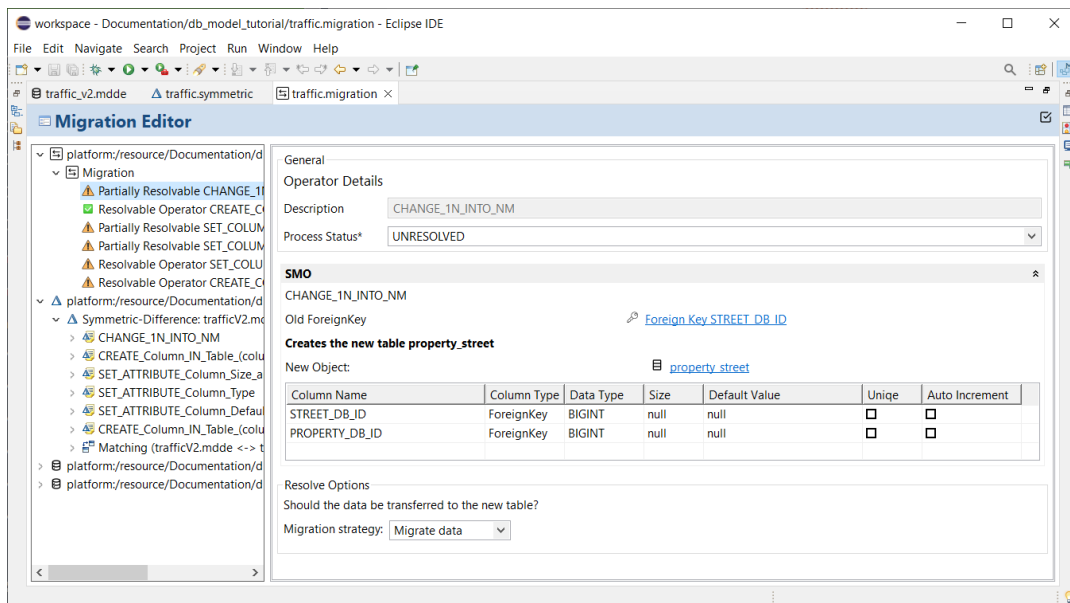


Figure 30: Partially Resolvable Operator

All operators requiring data migration usually belong to the group *Breaking & Resolvable*. For these operators, it is necessary to specify how and whether EcolveDB should migrate the

data. For this purpose, EvolveDB provides different migration strategies. These strategies depend on the change operator. To make this more precise, we will look at three examples below.

**CHANGE\_1N\_INTO\_NM:** Changing the multiplicity between property and street requires several operations (fig. 31). First, we must create a new table with the two foreign keys. Second, existing data should be transferred to the new table. Finally, we can delete the old foreign key. EvolveDB allows us to choose between two different migration strategies. We can choose a migration strategy by selecting a value from the dropdown box (fig. 31). In this case, the possibilities are simple because we only have two options. If *Migrate Data* is selected, the data will be transferred to the new table. Otherwise, the data will not be migrated and could be lost.

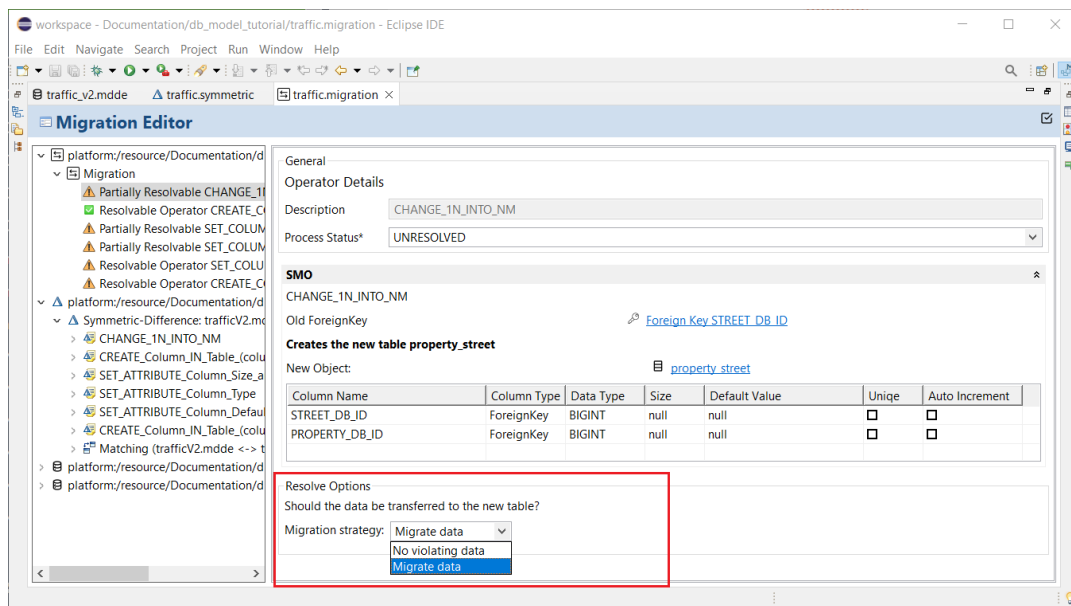


Figure 31: Migration strategy

**SET\_COLUMN\_SIZE:** In this example, we reduced the size of *VARCHAR* column from 255 to 40 (fig. 32). If the column contains values longer than the new size, the migration will abort with an error. The following enumeration lists all possible migration strategies:

- **No violating data:** no values invalidate the new constraint.
- If the column has a default value and has no unique constraint:
  1. **Set column to default value:** all values are set to the default value.
  2. **Set violating rows to the default value:** all values longer than 40 are set to the default value.
- If the column is nullable:
  1. **Set column to null:** all values are set to null.

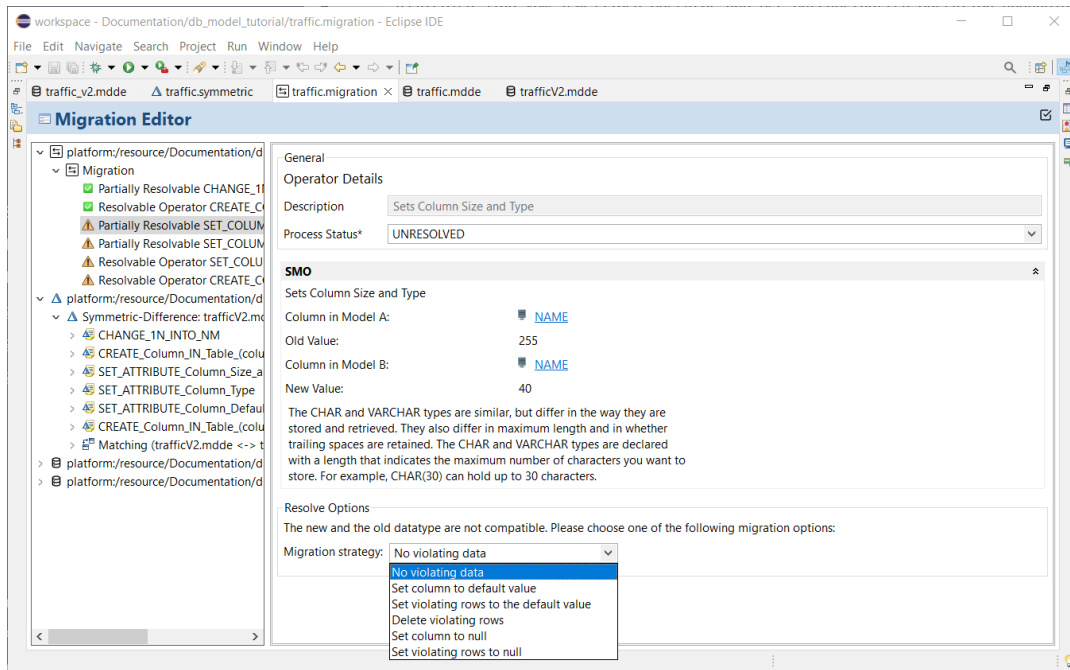


Figure 32: Migration strategy

## 2. Set violating rows to null: all values longer than 40 are set to null.

**SET\_COLUMN\_TYPE:** In the third example, we consider the type change. Type changes also belong to the second group (*Breaking & Resolvable*) because existing data could be incompatible with the target type. EvolveDB automatically checks if the old and the new types are compatible. Depending on the result, different migration strategies are offered. In our example shown (fig. 33), we changed the column type from *TIMESTAMP* to *DATETIME*. Because both data types are compatible, we do not need to choose a migration strategy in this case.

### 3.4.3 Not Automatically Resolvable

-TBD-

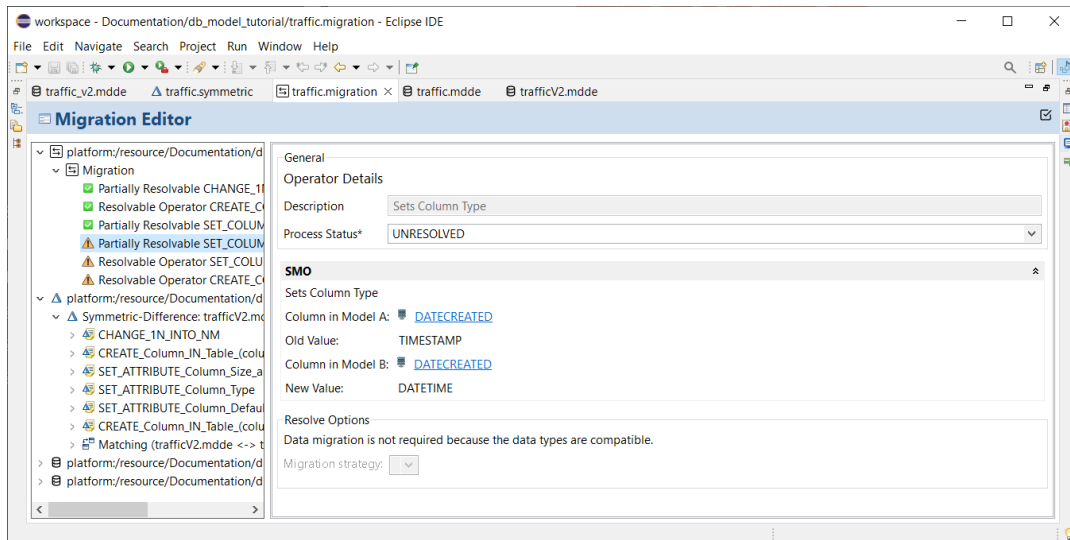


Figure 33: Migration strategy

### 3.5 Generate Migration Scripts

After choosing a migration strategy and setting the process status, we continue generating the migration script. Select the migration model in Package Explorer and open the context menu with a right-click. Choose **MDSE** ⇒ **Generate SQL Migrations** (fig. 34). We can also start the generation via the toolbar in the Migration Editor.

A wizard dialog opens. We have to choose the generator on the first page of the dialog. Since the SQL commands can differ slightly depending on the underlying database or the database version, EvolveDB offers an extension point for generators. In the default installation, only the MySQL-compatible generator is included (fig. 35).

On the second page of the dialog, we must select the location for the SQL script (fig. 36). Currently, we can only select a project in the workspace as the storage location. With *Finish* the generation process starts. If the file is not visible immediately afterwards, the project must be *refreshed*.

### 3 EvolveDB

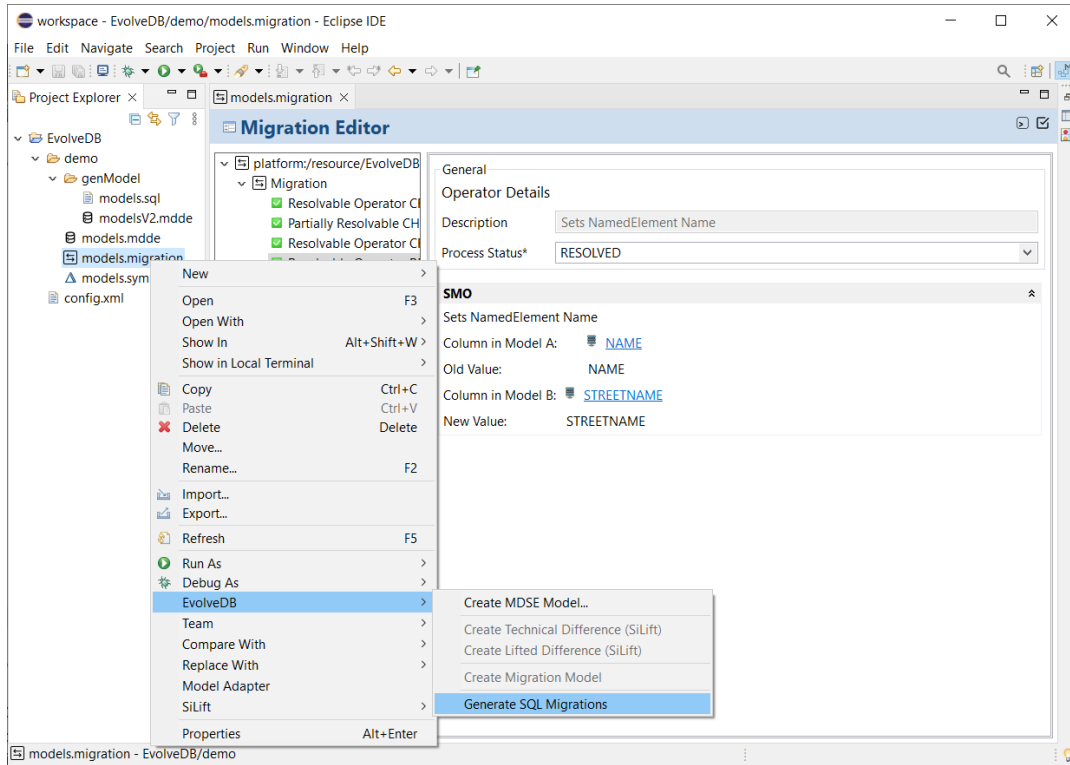


Figure 34: Create SQL migration script

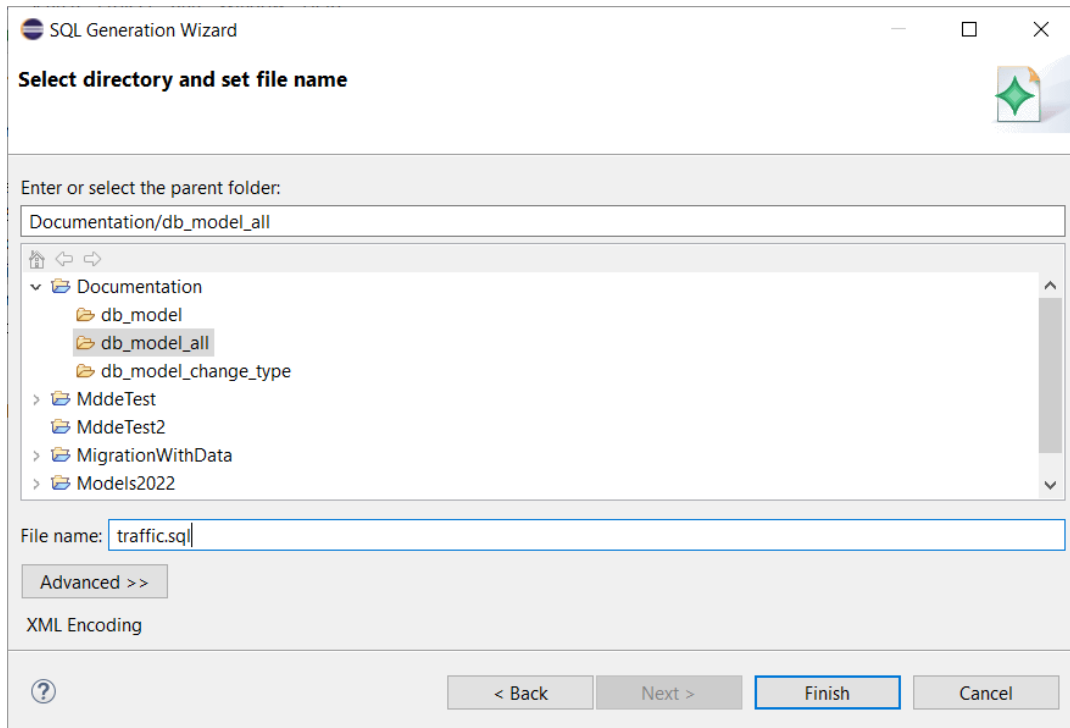


Figure 36: Choose file location

## 4 Migration Scripts

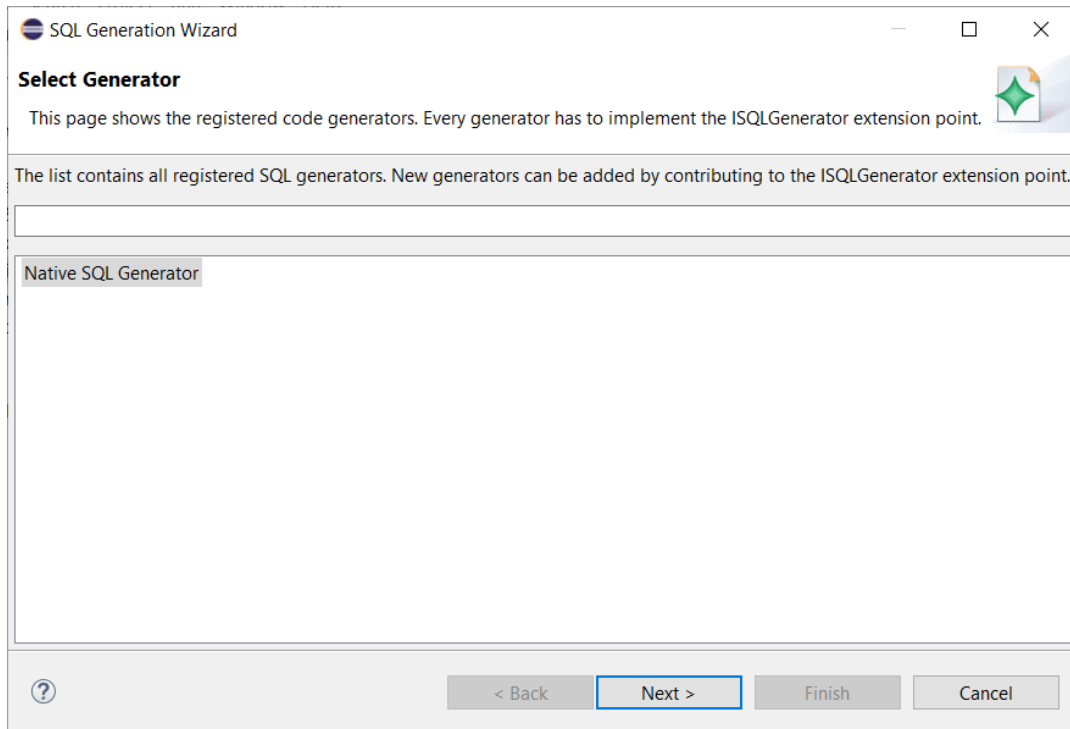


Figure 35: Select generator

## 4 Migration Scripts

In this chapter we will have a look at the generated migration script. Listing 1 shows the generated script when we use the model from our running example (fig. 27).

```
USE marburg_2020_models3; 1
START TRANSACTION; 2
3
-- Creates an history table for deleted and updated values 4
CREATE TABLE IF NOT EXISTS 'marburg_2020_models3'.' 5
  mdde_history' (
  'DB_ID' BIGINT NOT NULL AUTO_INCREMENT, 6
  'TABLENAME' VARCHAR(255) NOT NULL, 7
  'COLUMN_DB_ID' BIGINT NOT NULL, 8
  'COLUMN_NAME' VARCHAR(255) NOT NULL, 9
  'VALUE' BLOB NULL, 10
  'CHANGEDATE' DATETIME NOT NULL, 11
  PRIMARY KEY ('DB_ID')); 12
13
-- Add the new column LONGITUDE in Table crossroad 14
ALTER TABLE 'crossroad' 15
ADD COLUMN 'LONGITUDE' BIGINT; 16
```

#### 4 Migration Scripts

```
-- Change default value of lanes 17
ALTER TABLE 'crossroad' CHANGE COLUMN 'LANES' 'LANES' BIGINT 18
    NULL DEFAULT 1 ; 19

-- Add the new column LATITUDE in Table crossroad 20
ALTER TABLE 'crossroad' 21
ADD COLUMN 'LATITUDE' BIGINT; 22
23
24

-- Change column type of datecreated 25
ALTER TABLE 'street' CHANGE COLUMN 'DATECREATED' ' 26
    DATECREATED' DATETIME(0) NULL ; 27

-- Create Table property_street 28
CREATE TABLE IF NOT EXISTS property_street ( 29
'STREET_DB_ID' BIGINT NOT NULL, 30
'PROPERTY_DB_ID' BIGINT NOT NULL 31
,PRIMARY KEY('STREET_DB_ID' , 'PROPERTY_DB_ID' ), 32
CONSTRAINT 'property_crossroad_ibfk2' 33
FOREIGN KEY ('STREET_DB_ID') 34
REFERENCES 'street' ('DB_ID') 35
ON DELETE RESTRICT 36
ON UPDATE RESTRICT, 37
CONSTRAINT 'property_crossroad_ibfk3' 38
FOREIGN KEY ('PROPERTY_DB_ID') 39
REFERENCES 'property' ('DB_ID') 40
ON DELETE RESTRICT 41
ON UPDATE RESTRICT 42
); 43
44

BEGIN; 45
SET @safe_mode = @@SQL_SAFE_UPDATES; 46
SET SQL_SAFE_UPDATES = 0; 47
48

-- Migrate data to the new table 49
INSERT INTO 'property_street' (PROPERTY_DB_ID , STREET_DB_ID 50
)
SELECT DB_ID, STREET_DB_ID FROM property WHERE STREET_DB_ID 51
    IS NOT NULL; 52

SET SQL_SAFE_UPDATES = @safe_mode; 53
COMMIT; 54
-- If executing the script fails, we suggest a rollback. 55
56
```



## 4 Migration Scripts

```
-- Drop foreign key in property 57
ALTER TABLE 'property' DROP FOREIGN KEY 'ibfk_street'; 58
ALTER TABLE 'property' DROP COLUMN 'STREET_DB_ID'; 59
60
-- Find violating rows 61
SET @sql_mode = @@SESSION.sql_mode; 62
set @@SESSION.sql_mode = ''; 63
DROP TEMPORARY TABLE IF EXISTS my_temp_id_table; 64
CREATE TEMPORARY TABLE my_temp_id_table 65
    SELECT DB_ID from street v where LENGTH('NAME') > 40; 66
set @@SESSION.sql_mode = @sql_mode; 67
68
BEGIN; 69
SET @safe_mode = @@SQL_SAFE_UPDATES; 70
SET SQL_SAFE_UPDATES = 0; 71
72
-- Insert violating values into the history table 73
INSERT INTO 'marburg_2020_models3'.'mdde_history' 74
('TABLENAME', 75
'COLUMN_DB_ID', 76
'COLUMN_NAME', 77
'VALUE', 78
'CHANGEDATE') 79
SELECT 'street', DB_ID, 'NAME', NAME, now() from street s 80
    where LENGTH('NAME') > 40; 81
82
-- Set violating rows to the default value 82
UPDATE 'street' SET 'NAME' = 'unnamed_road' where DB_ID in ( 83
    Select DB_ID from my_temp_id_table); 84
85
SET SQL_SAFE_UPDATES = @safe_mode; 85
COMMIT; 86
-- If executing the script fails, we suggest a rollback. 87
88
DROP TEMPORARY TABLE IF EXISTS my_temp_id_table; 89
90
-- Change column size of name 91
ALTER TABLE 'street' CHANGE COLUMN 'NAME' 'NAME' VARCHAR(40) 92
    NULL DEFAULT 'unnamed_road'; 93
COMMIT; 93
-- If executing the script fails, we suggest a rollback. 94
```

Listing 1: Generated SQL code

If we apply the script to the original schema in figure 3, we obtain the migrated schema

#### 4 Migration Scripts

shown in figure 37. The new columns *LATITUDE* and *LONGITUDE* have been added. The original foreign key relationship between the *property* and *street* tables was moved to the new cross-reference table *property\_street*. Furthermore, the type of the *DATECREATED* column was changed to *datetime* and the size of the *NAME* column was reduced to 40.

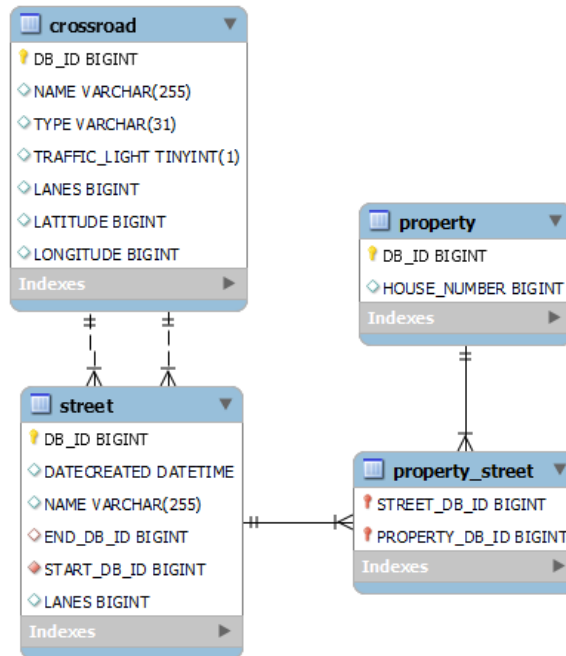


Figure 37: Das migrierte Datenbankschema

## 4 Migration Scripts

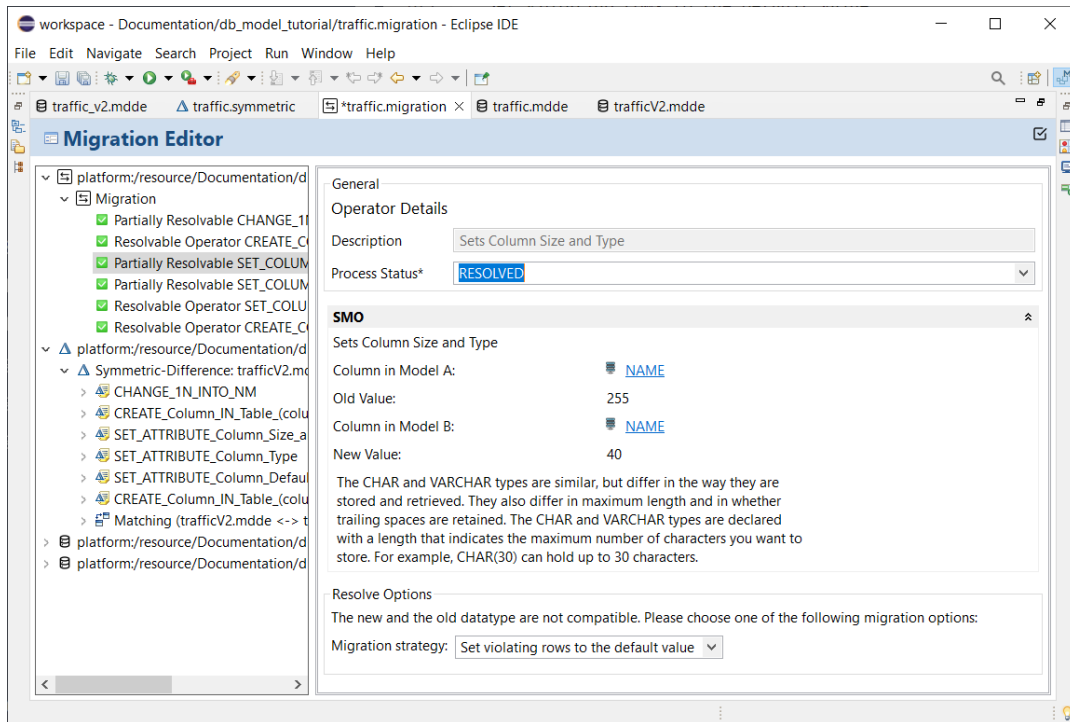


Figure 38: Migrationsstrategie mit History Tabelle

### 4.1 History Table

As mentioned earlier, EvolveDB provides default migration options. For example, we migrated the data from the old foreign key into the new table. In our running example, we also considered a case 3.4.2 where we had to select a migration strategy. To avoid data loss, EvolveDB creates a history table. Any values that are overwritten by one of the migration options are inserted into the history table. In our example, we chose the option to set all violating values to the default value 'unnamed road' (fig. 38).

```
-- Creates an history table for deleted and updated values 1
CREATE TABLE IF NOT EXISTS 'marburg_2020_models3'.' 2
  mdde_history' ( 3
  'DB_ID' BIGINT NOT NULL AUTO_INCREMENT, 4
  'TABLENAME' VARCHAR(255) NOT NULL, 5
  'COLUMN_DB_ID' BIGINT NOT NULL, 6
  'COLUMN_NAME' VARCHAR(255) NOT NULL, 7
  'VALUE' BLOB NULL, 8
  'CHANGEDATE' DATETIME NOT NULL, 9
  PRIMARY KEY ('DB_ID')); 10
  11
-- Create temporary table 12
SET @sql_mode = @@SESSION.sql_mode; 13
```

## 4 Migration Scripts

```
set @@SESSION.sql_mode = '';                                14
DROP TEMPORARY TABLE IF EXISTS my_temp_id_table;         15
CREATE TEMPORARY TABLE my_temp_id_table                  16
    SELECT DB_ID from street v where LENGTH('NAME') > 40; 17
set @@SESSION.sql_mode = @sql_mode;                       18
                                                            19
-- Find violating rows                                     20
BEGIN;                                                     21
SET @safe_mode = @@SQL_SAFE_UPDATES;                     22
SET SQL_SAFE_UPDATES = 0;                                  23
INSERT INTO 'marburg_2020_models3'.'mdde_history'        24
('TABLENAME',                                           25
 'COLUMN_DB_ID',                                         26
 'COLUMN_NAME',                                          27
 'VALUE',                                                28
 'CHANGEDATE')                                          29
SELECT 'street', DB_ID, 'NAME', NAME, now() from street s 30
    where LENGTH('NAME') > 40;                             31
                                                            32
-- Set violating rows to the default value                32
UPDATE 'street' SET 'NAME' = 'unnamed_road' where DB_ID in ( 33
    Select DB_ID from my_temp_id_table);
SET SQL_SAFE_UPDATES = @safe_mode;                       34
COMMIT;                                                    35
                                                            36
-- If executing the script fails, we suggest a rollback. 37
DROP TEMPORARY TABLE IF EXISTS my_temp_id_table;        38
                                                            39
-- Change column type and size of name                    40
ALTER TABLE 'street' CHANGE COLUMN 'NAME' 'NAME' VARCHAR(40) 41
    NULL DEFAULT 'unnamed_road';
```

Listing 2: Generated SQL code

The history table (see Fig. 39) is created between line three and line twelve. Subsequently, all columns whose contents are longer than 40 are transferred to the history table. Table 2 shows the history table after we performed the migration. One value was longer than 40 and was replaced by the default value *unnamed road*. Table 3 shows the content of table *street* after migration.

#### 4 Migration Scripts

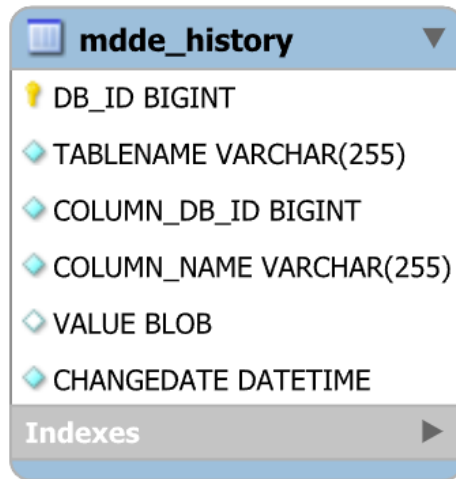


Figure 39: History table

DB_ID	TABLENAME	COL..._ID	COL..._NAME	VALUE	CHANGEDATE
1	street	2	NAME	BLOB	2022-05-20 15:41:57

Table 2: History table after executing the migration script.

DB_ID	DATECREATED	NAME	END_DB_ID	START_DB_ID	LANES
1	1976-05-20 15:41:57	Bahnhofstrasse	2	1	2
2	2019-08-12 13:25:57	unnamed road	2	3	3
3	1999-10-02 17:05:41	Seltersweg	3	1	1

Table 3: Street table after executing the migration script.